

k-Nearest Neighbors: Theory and Practice

September 10, 2025

1 Introduction

k-Nearest Neighbors (k-NN) is a non-parametric, instance-based learner: predictions are made by looking up the closest training samples. It is simple and often competitive on low-dimensional, well-scaled data, but can be sensitive to feature scaling and suffers in high dimensions.

2 Theory and Formulas

Given a query \mathbf{x} , find its k nearest neighbors under a distance metric $d(\cdot, \cdot)$ (e.g., Euclidean, Manhattan). For classification, predict by majority vote among the labels of these neighbors; with distance weighting (**weights=distance**), closer neighbors have higher influence. For regression, predict the average (or distance-weighted average) of neighbor targets.

Computationally, naive search costs $\mathcal{O}(nd)$ per query with n samples and d features. Tree-based indices (KDTree/BallTree) can accelerate queries in moderate dimensions. k-NN is affected by the curse of dimensionality; proper feature scaling and metric choice are critical.

3 Applications and Tips

- **Choose k:** tune via cross-validation; odd k helps avoid ties in binary classification.
- **Scaling:** standardize features or use pipelines; scale-sensitive.
- **Metric:** try Euclidean vs Manhattan; consider domain-specific distances.
- **Weights:** **uniform** vs **distance**; weighting can help with class overlap.
- **Complexity:** prediction cost grows with data size; consider approximate neighbors for large datasets.

4 Python Practice

Run the script in this chapter directory to generate figures into **figures/**.

Listing 1: Generate k-NN figures

```
1 python gen_knn_figures.py
```

Listing 2: gen_knn_figures.py

```

1  """
2  Figure generator for the k-NN chapter.
3
4  Generates illustrative figures and saves them into the chapter's 'figures/'
5  folder next to this script, regardless of current working directory.
6
7  Requirements:
8  - Python 3.8+
9  - numpy, matplotlib, scikit-learn
10
11  Install (if needed):
12      pip install numpy matplotlib scikit-learn
13
14  This script avoids newer or experimental APIs for broader compatibility.
15  """
16  from __future__ import annotations
17
18  import os
19  import numpy as np
20  import matplotlib.pyplot as plt
21  from matplotlib.colors import ListedColormap
22
23  try:
24      from sklearn.datasets import make_moons, make_regression,
25          make_classification
26      from sklearn.neighbors import KNeighborsClassifier, KNeighborsRegressor
27      from sklearn.preprocessing import StandardScaler
28      from sklearn.pipeline import make_pipeline
29      from sklearn.model_selection import cross_val_score
30  except Exception:
31      raise SystemExit(
32          "Missing scikit-learn. Please install with: pip install scikit-learn"
33      )
34
35  def _ensure_figures_dir(path: str | None = None) -> str:
36      """Create figures directory under this chapter regardless of CWD."""
37      if path is None:
38          base = os.path.dirname(os.path.abspath(__file__))
39          path = os.path.join(base, "figures")
40      os.makedirs(path, exist_ok=True)
41      return path
42
43
44  def _plot_decision_boundary(ax, clf, X, y, title: str):
45      x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
46      y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
47      xx, yy = np.meshgrid(
48          np.linspace(x_min, x_max, 400), np.linspace(y_min, y_max, 400)
49      )
50      Z = clf.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
51      cmap_light = ListedColormap(["#FFEEEE", "#EEEEFF"])

```

```

52     cmap_bold = ListedColormap(["#E74C3C", "#3498DB"])
53     ax.contourf(xx, yy, Z, cmap=cmap_light, alpha=0.8, levels=np.unique(Z).
        size)
54     ax.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolors="k", s=20)
55     ax.set_title(title)
56     ax.set_xlabel("Feature 1")
57     ax.set_ylabel("Feature 2")
58
59
60 def fig_knn_k_compare(out_dir: str) -> str:
61     np.random.seed(0)
62     X, y = make_moons(n_samples=500, noise=0.3, random_state=0)
63     models = [
64         (KNeighborsClassifier(n_neighbors=1), "k=1 (high variance)"),
65         (KNeighborsClassifier(n_neighbors=15), "k=15 (smoother)")
66     ]
67     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
        sharey=True)
68     for ax, (m, title) in zip(axes, models):
69         m.fit(X, y)
70         _plot_decision_boundary(ax, m, X, y, f"k-NN: {title}")
71     fig.suptitle("Effect of k on decision boundary")
72     out_path = os.path.join(out_dir, "knn_k_compare.png")
73     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
74     fig.savefig(out_path)
75     plt.close(fig)
76     return out_path
77
78
79 def fig_knn_metric_compare(out_dir: str) -> str:
80     np.random.seed(1)
81     X, y = make_moons(n_samples=500, noise=0.28, random_state=1)
82     models = [
83         (KNeighborsClassifier(n_neighbors=11, metric="euclidean"), "metric=
            euclidean"),
84         (KNeighborsClassifier(n_neighbors=11, metric="manhattan"), "metric=
            manhattan"),
85     ]
86     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
        sharey=True)
87     for ax, (m, title) in zip(axes, models):
88         m.fit(X, y)
89         _plot_decision_boundary(ax, m, X, y, f"k-NN: {title}")
90     fig.suptitle("Effect of distance metric")
91     out_path = os.path.join(out_dir, "knn_metric_compare.png")
92     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
93     fig.savefig(out_path)
94     plt.close(fig)
95     return out_path
96
97
98 def fig_knn_scaling_effect(out_dir: str) -> str:
99     np.random.seed(2)
100    X, y = make_classification(

```

```

101         n_samples=600,
102         n_features=2,
103         n_informative=2,
104         n_redundant=0,
105         n_clusters_per_class=1,
106         class_sep=1.0,
107         random_state=2,
108     )
109     # Impose different scales on features
110     X_scaled_variance = X.copy()
111     X_scaled_variance[:, 0] *= 8.0 # make feature 0 dominate distances
112
113     knn_raw = KNeighborsClassifier(n_neighbors=11)
114     knn_std = make_pipeline(StandardScaler(), KNeighborsClassifier(n_neighbors
        =11))
115
116     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
        sharey=True)
117     knn_raw.fit(X_scaled_variance, y)
118     _plot_decision_boundary(axes[0], knn_raw, X_scaled_variance, y, "Without
        scaling")
119     knn_std.fit(X_scaled_variance, y)
120     _plot_decision_boundary(axes[1], knn_std, X_scaled_variance, y, "With
        StandardScaler")
121     fig.suptitle("Feature scaling impact on k-NN")
122     out_path = os.path.join(out_dir, "knn_scaling_effect.png")
123     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
124     fig.savefig(out_path)
125     plt.close(fig)
126     return out_path
127
128
129 def fig_knn_weight_compare(out_dir: str) -> str:
130     np.random.seed(3)
131     X, y = make_moons(n_samples=500, noise=0.32, random_state=3)
132     models = [
133         (KNeighborsClassifier(n_neighbors=11, weights="uniform"), "weights=
            uniform"),
134         (KNeighborsClassifier(n_neighbors=11, weights="distance"), "weights=
            distance"),
135     ]
136     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
        sharey=True)
137     for ax, (m, title) in zip(axes, models):
138         m.fit(X, y)
139         _plot_decision_boundary(ax, m, X, y, f"k-NN: {title}")
140     fig.suptitle("Uniform vs distance weighting")
141     out_path = os.path.join(out_dir, "knn_weight_compare.png")
142     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
143     fig.savefig(out_path)
144     plt.close(fig)
145     return out_path
146
147

```

```

148 def fig_knn_regression_curve(out_dir: str) -> str:
149     rng = np.random.RandomState(4)
150     # 1D regression:  $y = \sin(x) + \text{noise}$ 
151     X = np.sort(rng.uniform(-3.0, 3.0, size=150)).reshape(-1, 1)
152     y = np.sin(X).ravel() + rng.normal(scale=0.25, size=X.shape[0])
153
154     grid = np.linspace(-3.5, 3.5, 600).reshape(-1, 1)
155     models = [
156         (KNeighborsRegressor(n_neighbors=1), "k=1"),
157         (KNeighborsRegressor(n_neighbors=15), "k=15"),
158         (KNeighborsRegressor(n_neighbors=45), "k=45"),
159     ]
160     fig, ax = plt.subplots(figsize=(7.5, 4.2), dpi=160)
161     ax.scatter(X[:, 0], y, s=18, c="#555", alpha=0.7, label="data")
162     colors = ["#E74C3C", "#3498DB", "#2ECC71"]
163     for (m, title), col in zip(models, colors):
164         m.fit(X, y)
165         y_pred = m.predict(grid)
166         ax.plot(grid[:, 0], y_pred, color=col, lw=2, label=title)
167     ax.set_title("k-NN regression: smoothing vs k")
168     ax.set_xlabel("x")
169     ax.set_ylabel("y")
170     ax.legend()
171     ax.grid(True, linestyle=":", alpha=0.4)
172     out_path = os.path.join(out_dir, "knn_regression_curve.png")
173     fig.tight_layout()
174     fig.savefig(out_path)
175     plt.close(fig)
176     return out_path
177
178
179 def main():
180     out_dir = _ensure_figures_dir(None)
181     generators = [
182         fig_knn_k_compare,
183         fig_knn_metric_compare,
184         fig_knn_scaling_effect,
185         fig_knn_weight_compare,
186         fig_knn_regression_curve,
187     ]
188     print("Generating figures into:", os.path.abspath(out_dir))
189     for gen in generators:
190         try:
191             p = gen(out_dir)
192             print("Saved:", p)
193         except Exception as e:
194             print("Failed generating", gen.__name__, ":", e)
195
196
197 if __name__ == "__main__":
198     main()

```

5 Result

Effect of k on decision boundary




Figure 1: Decision boundaries for different k (1 vs 15).

Effect of distance metric




Figure 2: Effect of metric: Euclidean vs Manhattan.

Feature scaling impact on k-NN




Figure 3: Impact of feature scaling on decision boundary.

Uniform vs distance weighting




Figure 4: Uniform vs distance weighting in k-NN.

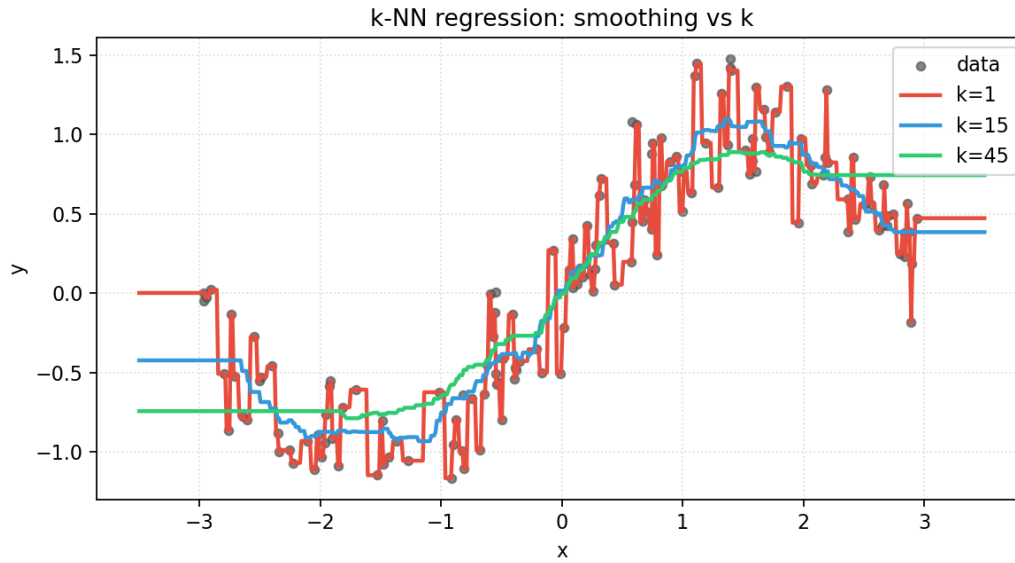


Figure 5: k-NN regression: smoothing effect as k increases.

6 Summary

k-NN is a simple yet powerful baseline for both classification and regression when features are well-scaled and dimensionality is moderate. Select k , metric, and weighting via validation, and use scaling to ensure distances are meaningful.