

Support Vector Machines: Theory and Practice

September 10, 2025

1 Introduction

Support Vector Machines (SVMs) are margin-based learners that find a decision boundary maximizing the margin between classes. With kernels, they represent complex non-linear decision boundaries while remaining convex to optimize.

2 Theory and Formulas

For linear, soft-margin SVM in primal form, given labeled data $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$ with $y_i \in \{-1, +1\}$:

$$\min_{\mathbf{w}, b, \xi \geq 0} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \quad (1)$$

$$\text{s.t.} \quad y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, n. \quad (2)$$

In the dual, data appear only via inner products; replacing them with kernels $K(\mathbf{x}, \mathbf{x}')$ provides non-linear SVMs. The decision function is

$$f(\mathbf{x}) = \sum_{i \in SV} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b, \quad \hat{y} = \text{sign } f(\mathbf{x}), \quad (3)$$

where SV are support vectors with non-zero multipliers α_i . RBF kernel $K(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$ is a common default. Hyperparameters C (slack penalty) and γ (kernel width) control regularization and boundary complexity.

3 Applications and Tips

- **Scaling:** always scale features; SVMs are sensitive to feature scales.
- **Hyperparameters:** tune C and γ (for RBF) via cross-validation; start with $C \in [0.1, 10]$, $\gamma \in [10^{-2}, 10]$ (log grid).
- **Class imbalance:** use `class_weight=balanced` to reweight classes.
- **Probability:** SVC supports probability with `probability=True` (adds calibration cost); otherwise use `decision_function`.
- **Multiclass:** SVC uses one-vs-one internally; `LinearSVC` uses one-vs-rest.

4 Python Practice

Run the script in this chapter directory to generate figures into `figures/`.

Listing 1: Generate SVM figures

```
1 python gen_svm_figures.py
```

Listing 2: `gen_svm_figures.py`

```
1 """
2 Figure generator for the SVM chapter.
3
4 Generates illustrative figures and saves them into the chapter's 'figures/'
5 folder next to this script, regardless of current working directory.
6
7 Requirements:
8 - Python 3.8+
9 - numpy, matplotlib, scikit-learn
10
11 Install (if needed):
12     pip install numpy matplotlib scikit-learn
13
14 This script avoids newer or experimental APIs for broader compatibility.
15 """
16 from __future__ import annotations
17
18 import os
19 import numpy as np
20 import matplotlib.pyplot as plt
21 from matplotlib.colors import ListedColormap
22
23 try:
24     from sklearn.datasets import make_moons, make_classification
25     from sklearn.svm import SVC
26 except Exception:
27     raise SystemExit(
28         "Missing scikit-learn. Please install with: pip install scikit-learn"
29     )
30
31
32 def _ensure_figures_dir(path: str | None = None) -> str:
33     """Create figures directory under this chapter regardless of CWD."""
34     if path is None:
35         base = os.path.dirname(os.path.abspath(__file__))
36         path = os.path.join(base, "figures")
37     os.makedirs(path, exist_ok=True)
38     return path
39
40
41 def _plot_decision_boundary(ax, clf, X, y, title: str):
42     x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
43     y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
44     xx, yy = np.meshgrid(
45         np.linspace(x_min, x_max, 400), np.linspace(y_min, y_max, 400)
```

```

46     )
47     Z = clf.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
48     cmap_light = ListedColormap(["#FFEEEE", "#EEEEFF"])
49     cmap_bold = ListedColormap(["#E74C3C", "#3498DB"])
50     ax.contourf(xx, yy, Z, cmap=cmap_light, alpha=0.8, levels=np.unique(Z).
        size)
51     ax.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolors="k", s=20)
52     ax.set_title(title)
53     ax.set_xlabel("Feature 1")
54     ax.set_ylabel("Feature 2")
55
56
57 def fig_svm_linear_vs_rbf(out_dir: str) -> str:
58     np.random.seed(0)
59     X, y = make_moons(n_samples=500, noise=0.3, random_state=0)
60     models = [
61         (SVC(kernel="linear", C=1.0, random_state=0), "Linear kernel"),
62         (SVC(kernel="rbf", C=1.0, gamma=1.0, random_state=0), "RBF kernel"),
63     ]
64     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
        sharey=True)
65     for ax, (m, title) in zip(axes, models):
66         m.fit(X, y)
67         _plot_decision_boundary(ax, m, X, y, f"SVM: {title}")
68     fig.suptitle("SVM: Linear vs RBF kernel")
69     out_path = os.path.join(out_dir, "svm_linear_vs_rbf.png")
70     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
71     fig.savefig(out_path)
72     plt.close(fig)
73     return out_path
74
75
76 def fig_svm_C_compare(out_dir: str) -> str:
77     np.random.seed(1)
78     X, y = make_moons(n_samples=500, noise=0.3, random_state=1)
79     models = [
80         (SVC(kernel="rbf", C=0.3, gamma=1.0, random_state=1), "C=0.3 (more
            regularized)"),
81         (SVC(kernel="rbf", C=100.0, gamma=1.0, random_state=1), "C=100 (less
            regularized)"),
82     ]
83     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
        sharey=True)
84     for ax, (m, title) in zip(axes, models):
85         m.fit(X, y)
86         _plot_decision_boundary(ax, m, X, y, f"SVM (RBF): {title}")
87     fig.suptitle("Effect of C (soft-margin)")
88     out_path = os.path.join(out_dir, "svm_C_compare.png")
89     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
90     fig.savefig(out_path)
91     plt.close(fig)
92     return out_path
93
94

```

```

95 def fig_svm_gamma_compare(out_dir: str) -> str:
96     np.random.seed(2)
97     X, y = make_moons(n_samples=500, noise=0.3, random_state=2)
98     models = [
99         (SVC(kernel="rbf", C=3.0, gamma=0.2, random_state=2), "gamma=0.2 (
100             smoother)"),
101         (SVC(kernel="rbf", C=3.0, gamma=5.0, random_state=2), "gamma=5.0 (
102             wiggly)")
103     ]
104     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
105                             sharey=True)
106     for ax, (m, title) in zip(axes, models):
107         m.fit(X, y)
108         _plot_decision_boundary(ax, m, X, y, f"SVM (RBF): {title}")
109     fig.suptitle("Effect of gamma (RBF width)")
110     out_path = os.path.join(out_dir, "svm_gamma_compare.png")
111     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
112     fig.savefig(out_path)
113     plt.close(fig)
114     return out_path
115
116 def fig_svm_margin_support_vectors(out_dir: str) -> str:
117     # Linearly separable-like data for margin visualization
118     X, y = make_classification(
119         n_samples=200,
120         n_features=2,
121         n_redundant=0,
122         n_informative=2,
123         n_clusters_per_class=1,
124         class_sep=2.0,
125         random_state=3,
126     )
127     clf = SVC(kernel="linear", C=1e3, random_state=3)
128     clf.fit(X, y)
129
130     fig, ax = plt.subplots(figsize=(6.5, 4.8), dpi=160)
131     _plot_decision_boundary(ax, clf, X, y, "Linear SVM with margin and SVs")
132
133     # Plot the margin lines using  $w^T x + b = \pm 1$ 
134     w = clf.coef_[0]
135     b = clf.intercept_[0]
136     # Create a grid line in x for margin lines
137     x_vals = np.linspace(X[:, 0].min() - 0.5, X[:, 0].max() + 0.5, 200)
138     # For  $y = -(w_0 x + b - m) / w_1$  with  $m$  in  $\{0, 1, -1\}$ 
139     if abs(w[1]) > 1e-12:
140         for m in [0.0, 1.0, -1.0]:
141             y_vals = -(w[0] * x_vals + b - m) / w[1]
142             style = "k-" if m == 0 else "k--"
143             ax.plot(x_vals, y_vals, style, lw=1.2, alpha=0.9)
144
145     # Highlight support vectors
146     sv = clf.support_vectors_
147     ax.scatter(sv[:, 0], sv[:, 1], s=80, facecolors="none", edgecolors="#000",

```

```

146         linewidths=1.5, label="SV")
147     ax.legend(loc="best")
148     out_path = os.path.join(out_dir, "svm_margin_support_vectors.png")
149     fig.tight_layout()
150     fig.savefig(out_path)
151     plt.close(fig)
152     return out_path
153
154 def fig_svm_decision_function(out_dir: str) -> str:
155     np.random.seed(4)
156     X, y = make_moons(n_samples=400, noise=0.25, random_state=4)
157     clf = SVC(kernel="rbf", C=2.0, gamma=1.0, random_state=4)
158     clf.fit(X, y)
159
160     x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
161     y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
162     xx, yy = np.meshgrid(
163         np.linspace(x_min, x_max, 500), np.linspace(y_min, y_max, 500)
164     )
165     Z = clf.decision_function(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
166
167     fig, ax = plt.subplots(figsize=(6.5, 5.0), dpi=160)
168     # Filled regions by sign
169     ax.contourf(xx, yy, (Z > 0).astype(int), levels=2, cmap=ListedColormap(["#
170         FFFFFFFF", "#EEEEFF"]), alpha=0.8)
171     # Decision function contours for -1, 0, +1
172     CS = ax.contour(xx, yy, Z, levels=[-1.0, 0.0, 1.0], colors=["k", "k", "k"
173         ], linestyles=["--", "-", "--"], linewidths=1.2)
174     ax.clabel(CS, inline=True, fontsize=8, fmt={-1.0: "-1", 0.0: "0", 1.0: "+1
175         "})
176     # Data points and SVs
177     ax.scatter(X[:, 0], X[:, 1], c=y, cmap=ListedColormap(["#E74C3C", "#3498DB
178         "]), edgecolors="k", s=20)
179     sv = clf.support_vectors_
180     ax.scatter(sv[:, 0], sv[:, 1], s=80, facecolors="none", edgecolors="#000",
181         linewidths=1.5, label="SV")
182     ax.set_title("RBF SVM: decision function and margins")
183     ax.set_xlabel("Feature 1")
184     ax.set_ylabel("Feature 2")
185     ax.legend(loc="best")
186     out_path = os.path.join(out_dir, "svm_decision_function.png")
187     fig.tight_layout()
188     fig.savefig(out_path)
189     plt.close(fig)
190     return out_path
191
192 def main():
193     out_dir = _ensure_figures_dir(None)
194     generators = [
195         fig_svm_linear_vs_rbf,
196         fig_svm_C_compare,
197         fig_svm_gamma_compare,

```

```

194     fig_svm_margin_support_vectors,
195     fig_svm_decision_function,
196 ]
197 print("Generating figures into:", os.path.abspath(out_dir))
198 for gen in generators:
199     try:
200         p = gen(out_dir)
201         print("Saved:", p)
202     except Exception as e:
203         print("Failed generating", gen.__name__, ":", e)
204
205
206 if __name__ == "__main__":
207     main()

```

5 Result

SVM: Linear vs RBF kernel




Figure 1: SVM decision boundaries: linear vs RBF kernel.

Effect of C (soft-margin)




Figure 2: Effect of soft-margin parameter C (RBF kernel).

Effect of gamma (RBF width)

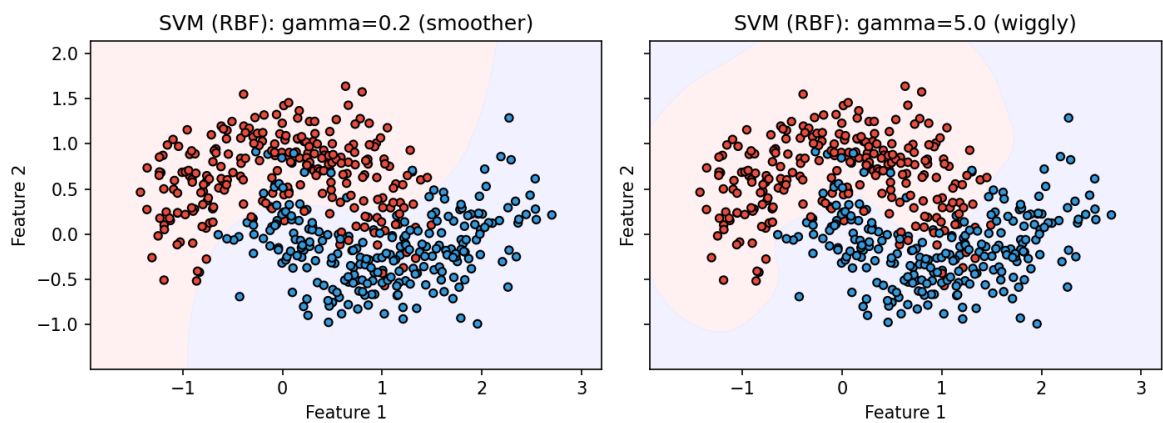


Figure 3: Effect of RBF gamma on boundary smoothness.

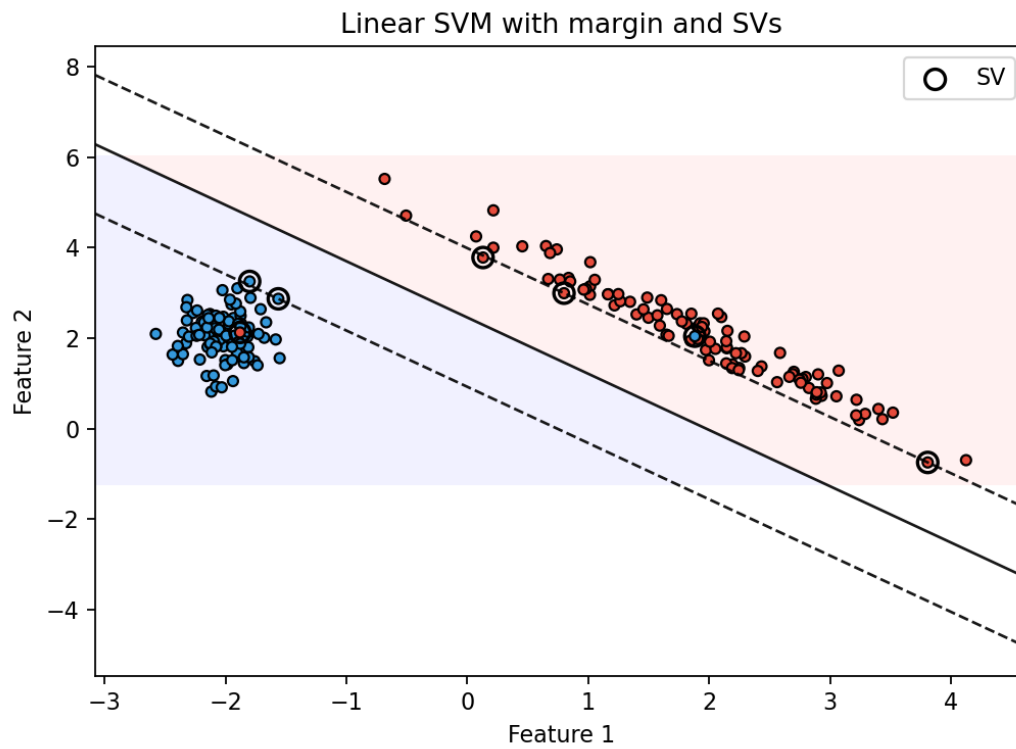


Figure 4: Linear SVM: margin lines and highlighted support vectors.




Figure 5: RBF SVM decision function: contours at -1, 0, +1.

6 Summary

SVMs maximize margins for robust classification and extend to non-linear problems via kernels. With proper scaling and tuning of C and kernel parameters, they deliver strong performance across many domains.