

XGBoost: Theory and Practice

September 10, 2025

1 Introduction

XGBoost is an efficient and scalable implementation of gradient boosted decision trees (GBDTs). It improves training speed, regularization, and accuracy through second-order optimization, tree sparsity-aware split finding, shrinkage, and subsampling.

2 Theory and Formulas

Gradient boosting fits an additive model $F_M(\mathbf{x}) = \sum_{m=1}^M f_m(\mathbf{x})$ of shallow trees by stage-wise optimization. XGBoost minimizes a regularized objective

$$\mathcal{L} = \sum_{i=1}^n \ell(y_i, \hat{y}_i) + \sum_{m=1}^M \Omega(f_m), \quad \Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2, \quad (1)$$

where T is the number of leaves and w are leaf scores. Using a second-order Taylor expansion of the loss around current predictions yields per-node sums of gradients g_i and Hessians h_i ; the split gain for left/right partitions L, R is

$$\text{Gain} = \frac{1}{2} \left(\frac{(\sum_{i \in L} g_i)^2}{\sum_{i \in L} h_i + \lambda} + \frac{(\sum_{i \in R} g_i)^2}{\sum_{i \in R} h_i + \lambda} - \frac{(\sum_{i \in L \cup R} g_i)^2}{\sum_{i \in L \cup R} h_i + \lambda} \right) - \gamma. \quad (2)$$

Regularization via λ, γ , shrinkage (learning rate), column/row subsampling, and maximum depth/leaf constraints control complexity and reduce overfitting.

3 Applications and Tips

- **Scaling and types:** handles numeric and one-hot encoded categorical features; no scaling required for trees.
- **Key hyperparameters:** `n_estimators`, `max_depth`, `learning_rate`, `subsample`, `colsample_bytree`, `reg_alpha/reg_lambda`.
- **Early stopping:** use validation with `eval_set` and `early_stopping_rounds`.
- **Imbalance:** set `scale_pos_weight` or use stratified sampling.
- **Interpretation:** start with built-in importances; prefer permutation or SHAP for robust insights.

4 Python Practice

Run the script in this chapter directory to generate figures into `figures/`.

Listing 1: Generate XGBoost figures

```
1 python gen_xgboost_figures.py
```

Listing 2: `gen_xgboost_figures.py`

```
1 """
2 Figure generator for the XGBoost chapter.
3
4 Generates illustrative figures and saves them into the chapter's 'figures/'
5 folder next to this script, regardless of current working directory.
6
7 Requirements:
8 - Python 3.8+
9 - numpy, matplotlib, scikit-learn
10 - xgboost (optional; falls back to scikit-learn GradientBoosting if missing)
11
12 Install (if needed):
13     pip install numpy matplotlib scikit-learn xgboost
14
15 This script avoids newer or experimental APIs for broader compatibility.
16 """
17 from __future__ import annotations
18
19 import os
20 import numpy as np
21 import matplotlib.pyplot as plt
22 from matplotlib.colors import ListedColormap
23
24 try:
25     import xgboost as xgb # type: ignore
26     HAS_XGB = True
27 except Exception:
28     xgb = None
29     HAS_XGB = False
30
31 from sklearn.datasets import make_moons, make_classification
32 from sklearn.model_selection import train_test_split
33 from sklearn.metrics import log_loss
34
35 try:
36     from sklearn.ensemble import GradientBoostingClassifier
37 except Exception:
38     GradientBoostingClassifier = None # type: ignore
39
40
41 def _ensure_figures_dir(path: str | None = None) -> str:
42     """Create figures directory under this chapter regardless of CWD."""
43     if path is None:
44         base = os.path.dirname(os.path.abspath(__file__))
45         path = os.path.join(base, "figures")
```

```

46     os.makedirs(path, exist_ok=True)
47     return path
48
49
50 def _plot_decision_boundary(ax, clf, X, y, title: str):
51     x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
52     y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
53     xx, yy = np.meshgrid(
54         np.linspace(x_min, x_max, 400), np.linspace(y_min, y_max, 400)
55     )
56     Z = clf.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
57     cmap_light = ListedColormap(["#FFEEEE", "#EEEEFF"])
58     cmap_bold = ListedColormap(["#E74C3C", "#3498DB"])
59     ax.contourf(xx, yy, Z, cmap=cmap_light, alpha=0.8, levels=np.unique(Z).
60                 size)
61     ax.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolors="k", s=20)
62     ax.set_title(title)
63     ax.set_xlabel("Feature 1")
64     ax.set_ylabel("Feature 2")
65
66 def _xgb_classifier(**kwargs):
67     if HAS_XGB:
68         params = dict(
69             n_estimators=200,
70             max_depth=3,
71             learning_rate=0.1,
72             subsample=1.0,
73             colsample_bytree=1.0,
74             objective="binary:logistic",
75             tree_method="hist",
76             random_state=0,
77             n_jobs=0,
78         )
79         params.update(kwargs)
80         return xgb.XGBClassifier(**params)
81     else:
82         if GradientBoostingClassifier is None:
83             raise RuntimeError("Neither xgboost nor GradientBoostingClassifier
84                                available.")
85         params = dict(
86             n_estimators=kwargs.get("n_estimators", 200),
87             max_depth=kwargs.get("max_depth", 3),
88             learning_rate=kwargs.get("learning_rate", 0.1),
89             random_state=0,
90         )
91         return GradientBoostingClassifier(**params)
92
93 def fig_xgb_decision_boundary_2class(out_dir: str) -> str:
94     np.random.seed(0)
95     X, y = make_moons(n_samples=500, noise=0.3, random_state=0)
96     clf = _xgb_classifier()
97     clf.fit(X, y)

```

```

98
99 fig, ax = plt.subplots(figsize=(5.5, 4.5), dpi=150)
100 title = "XGBoost boundary (max_depth=3, lr=0.1)" if HAS_XGB else "GBDT
      boundary (fallback)"
101 _plot_decision_boundary(ax, clf, X, y, title)
102 out_path = os.path.join(out_dir, "xgb_decision_boundary_2class.png")
103 fig.tight_layout()
104 fig.savefig(out_path)
105 plt.close(fig)
106 return out_path
107
108
109 def fig_xgb_learning_rate_compare(out_dir: str) -> str:
110     np.random.seed(1)
111     X, y = make_moons(n_samples=550, noise=0.3, random_state=1)
112     models = [
113         (_xgb_classifier(learning_rate=0.05, n_estimators=400), "learning_rate
              =0.05, n_estimators=400"),
114         (_xgb_classifier(learning_rate=0.3, n_estimators=150), "learning_rate
              =0.3, n_estimators=150"),
115     ]
116     fig, axes = plt.subplots(1, 2, figsize=(9.5, 4.2), dpi=150, sharex=True,
        sharey=True)
117     for ax, (m, title) in zip(axes, models):
118         m.fit(X, y)
119         label = ("XGBoost: " if HAS_XGB else "GBDT: ") + title
120         _plot_decision_boundary(ax, m, X, y, label)
121     fig.suptitle("Effect of learning_rate with trees budget")
122     out_path = os.path.join(out_dir, "xgb_learning_rate_compare.png")
123     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
124     fig.savefig(out_path)
125     plt.close(fig)
126     return out_path
127
128
129 def fig_xgb_max_depth_compare(out_dir: str) -> str:
130     np.random.seed(2)
131     X, y = make_moons(n_samples=600, noise=0.32, random_state=2)
132     models = [
133         (_xgb_classifier(max_depth=2, n_estimators=250), "max_depth=2"),
134         (_xgb_classifier(max_depth=4, n_estimators=250), "max_depth=4"),
135         (_xgb_classifier(max_depth=8, n_estimators=250), "max_depth=8"),
136     ]
137     fig, axes = plt.subplots(1, 3, figsize=(12.5, 4.2), dpi=150, sharex=True,
        sharey=True)
138     for ax, (m, title) in zip(axes, models):
139         m.fit(X, y)
140         label = ("XGBoost: " if HAS_XGB else "GBDT: ") + title
141         _plot_decision_boundary(ax, m, X, y, label)
142     fig.suptitle("Effect of max_depth")
143     out_path = os.path.join(out_dir, "xgb_max_depth_compare.png")
144     fig.tight_layout(rect=[0, 0.03, 1, 0.95])
145     fig.savefig(out_path)
146     plt.close(fig)

```

```

147     return out_path
148
149
150 def fig_xgb_feature_importances(out_dir: str) -> str:
151     X, y = make_classification(
152         n_samples=800,
153         n_features=10,
154         n_informative=4,
155         n_redundant=3,
156         n_repeated=0,
157         random_state=7,
158         shuffle=True,
159     )
160     clf = _xgb_classifier(n_estimators=300, max_depth=4, learning_rate=0.1)
161     clf.fit(X, y)
162     importances = getattr(clf, "feature_importances_", None)
163     if importances is None:
164         # Fallback: uniform zeros to avoid crash
165         importances = np.zeros(X.shape[1], dtype=float)
166
167     fig, ax = plt.subplots(figsize=(7.0, 4.0), dpi=160)
168     idx = np.arange(importances.size)
169     ax.bar(idx, importances, color="#F39C12")
170     ax.set_xticks(idx)
171     ax.set_xticklabels([f"f{i}" for i in idx])
172     ax.set_ylabel("importance")
173     title = "XGBoost feature importances" if HAS_XGB else "GBDT feature
        importances"
174     ax.set_title(title)
175     ax.set_ylim(0, max(0.25, float(importances.max()) + 0.05))
176     for i, v in enumerate(importances):
177         ax.text(i, v + 0.01, f"{v:.2f}", ha="center", va="bottom", fontsize=8)
178     out_path = os.path.join(out_dir, "xgb_feature_importances.png")
179     fig.tight_layout()
180     fig.savefig(out_path)
181     plt.close(fig)
182     return out_path
183
184
185 def fig_xgb_eval_logloss_curve(out_dir: str) -> str:
186     np.random.seed(3)
187     X, y = make_classification(
188         n_samples=1200,
189         n_features=15,
190         n_informative=5,
191         n_redundant=5,
192         random_state=3,
193     )
194     X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3,
        random_state=3)
195
196     if HAS_XGB:
197         clf = _xgb_classifier(n_estimators=300, learning_rate=0.1, max_depth
            =4)

```

```

198     clf.fit(X_train, y_train, eval_set=[(X_train, y_train), (X_val, y_val)
199           ], eval_metric="logloss", verbose=False)
200     res = clf.evals_result()
201     tr = np.array(res.get("validation_0", {}).get("logloss", []), dtype=
202         float)
203     va = np.array(res.get("validation_1", {}).get("logloss", []), dtype=
204         float)
205 else:
206     # Fallback using staged decision on GradientBoosting
207     clf = _xgb_classifier(n_estimators=300, learning_rate=0.1, max_depth
208         =3)
209     clf.fit(X_train, y_train)
210     tr_list, va_list = [], []
211     # GradientBoostingClassifier provides staged_predict_proba
212     if hasattr(clf, "staged_predict_proba"):
213         for y_tr_prob, y_va_prob in zip(clf.staged_predict_proba(X_train),
214             clf.staged_predict_proba(X_val)):
215             tr_list.append(log_loss(y_train, y_tr_prob))
216             va_list.append(log_loss(y_val, y_va_prob))
217     else:
218         # Last resort: single-point curves
219         y_tr_prob = clf.predict_proba(X_train)
220         y_va_prob = clf.predict_proba(X_val)
221         tr_list = [log_loss(y_train, y_tr_prob)]
222         va_list = [log_loss(y_val, y_va_prob)]
223     tr, va = np.array(tr_list), np.array(va_list)
224
225 fig, ax = plt.subplots(figsize=(6.8, 4.2), dpi=160)
226 ax.plot(np.arange(1, len(tr) + 1), tr, label="train logloss", color="#2
227     ECC71")
228 ax.plot(np.arange(1, len(va) + 1), va, label="valid logloss", color="#
229     E74C3C")
230 ax.set_xlabel("n_trees")
231 ax.set_ylabel("logloss")
232 ax.set_title("Evaluation curve (logloss vs trees)")
233 ax.legend()
234 ax.grid(True, linestyle=":", alpha=0.4)
235 out_path = os.path.join(out_dir, "xgb_eval_logloss_curve.png")
236 fig.tight_layout()
237 fig.savefig(out_path)
238 plt.close(fig)
239 return out_path
240
241 def main():
242     out_dir = _ensure_figures_dir(None)
243     generators = [
244         fig_xgb_decision_boundary_2class,
245         fig_xgb_learning_rate_compare,
246         fig_xgb_max_depth_compare,
247         fig_xgb_feature_importances,
248         fig_xgb_eval_logloss_curve,
249     ]
250     print("Generating figures into:", os.path.abspath(out_dir))

```

```

245     if not HAS_XGB:
246         print("xgboost not found; falling back to GradientBoostingClassifier
           where possible.")
247     for gen in generators:
248         try:
249             p = gen(out_dir)
250             print("Saved:", p)
251         except Exception as e:
252             print("Failed generating", gen.__name__, ":", e)
253
254
255 if __name__ == "__main__":
256     main()

```

5 Result




Figure 1: XGBoost decision boundary on a 2-class dataset.

Effect of learning_rate with trees budget

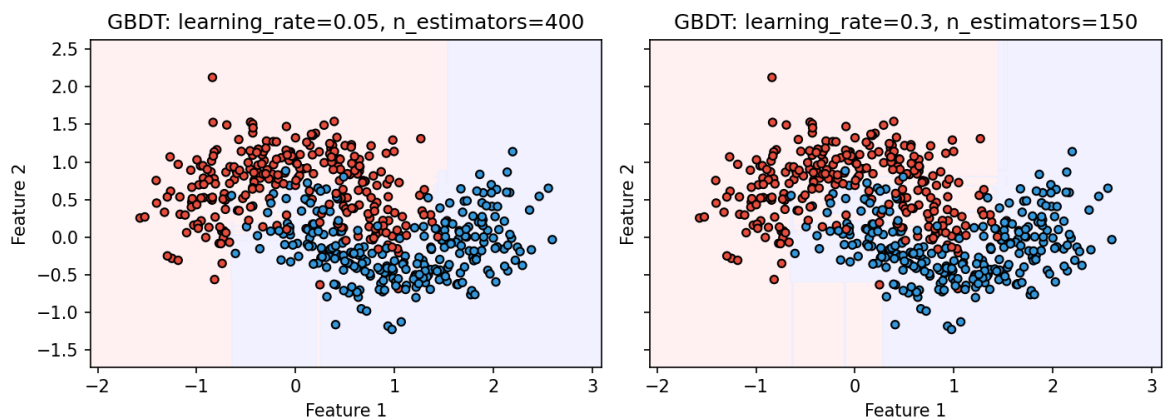


Figure 2: Learning rate effect with a budget of trees.

Effect of max_depth




Figure 3: Decision boundaries under different max_depth.




Figure 4: Feature importances from XGBoost.

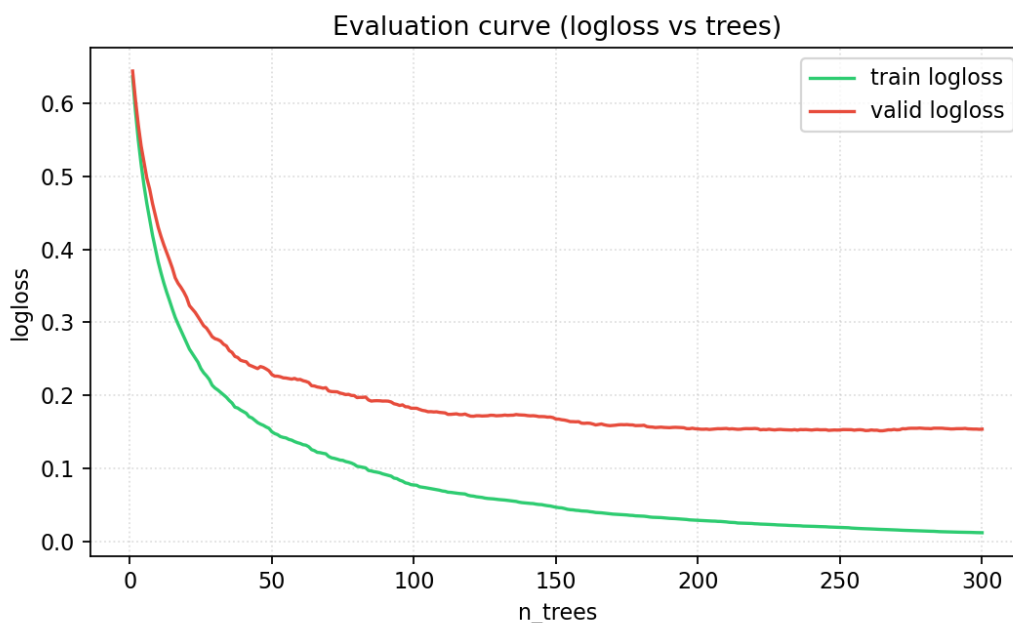


Figure 5: Training/validation logloss vs number of trees.

6 Summary

XGBoost combines efficient tree boosting with strong regularization and advanced split finding. With careful tuning of depth, learning rate, and sampling, it delivers state-of-the-art performance on many tabular tasks.